

Image Vectorizer with Curve Fitting:

A Computational-Geometry Pipeline from
Raster Pixels to Smooth Bézier Paths

Nantapong Ruangsuksriwong

Student ID 6632107621 GitHub: @teamangkorn

Computational Geometry, Term 2/2568

Department of Computer Engineering, Chulalongkorn University

Abstract

We present a complete raster-to-vector image converter built around three classical computational-geometry algorithms: Moore-Neighbour boundary tracing, Ramer–Douglas–Peucker (RDP) polyline simplification, and Schneider’s recursive cubic-Bézier fit (evaluated throughout with the de Casteljau subdivision scheme). We compare two *baselines* — a raw pixel-outline polygon emitter and an RDP-only polyline emitter — against the *proposed* pipeline that adds Schneider Bézier fitting on top of RDP. On synthetic inputs scaled across $n \in \{10^3, 10^4, 10^5\}$ traced boundary vertices, the proposed pipeline produces SVG output that is **20–27× smaller** than the raw baseline (e.g. 1754478 → 84630 bytes at $n = 10^5$, a 95.2% reduction) and **1.7–1.8× smaller** than the RDP-only baseline, while adding less than 1 ms of curve fitting overhead even on the largest input. The proposed pipeline is also marginally *faster* end-to-end (5% at $n = 10^5$) because RDP shrinks the data before SVG serialization. All three modes share the same Moore-Neighbour tracer, so the asymptotic cost is dominated by tracing and remains $O(WH)$ in image area.

Keywords: image vectorization, Bézier curves, polyline simplification, Moore-Neighbour tracing, computational geometry.

1 Introduction

A raster image is a regular grid of colour samples; a vector image describes the same content as resolution-independent geometric primitives. Vectorization is therefore the *inverse* of rasterization, and is in general ill-posed — the original geometry is not recoverable from samples. Practical vectorizers introduce strong priors (piecewise-smooth boundaries, sparse control polygons) and reformulate the task as one of *geometric approximation*. This is exactly the regime in which the algorithms taught in our Computational Geometry course apply.

In this project we build such a vectorizer end-to-end and use it to quantify the benefit of curve fitting: how much smaller (and how much slower, if at all) is a Bézier-based output than the simpler polygonal baselines that one could emit using only the first half of the pipeline? We answer that question with a controlled scaling experiment over three orders of magnitude in input size.

Contributions. The specific contributions of this work are:

- A complete raster-to-vector pipeline in < 1000 lines of C++17 with a single public-domain dependency (`stb_image.h`).
- A controlled three-way ablation experiment that isolates the cost and benefit of each pipeline stage independently.
- Experimental evidence that adding Schneider’s cubic-Bézier fit makes the pipeline *faster* end-to-

end — a counterintuitive result driven entirely by SVG serialization savings.

- A per-stage timing breakdown that focuses future optimisation effort on the dominating Moore-Neighbour tracing stage.

Paper organisation. Section 2 formalises the vectorization problem. Section 3 surveys related work and defines the baselines. Section 4 presents algorithm design and complexity analysis. Section 5 covers implementation. Section 6 describes correctness validation. Sections 7 and 8 report the experimental study. Sections 9–10 conclude. Appendix B derives the least-squares system at the core of Schneider’s fit.

2 Problem Definition

2.1 Background and Significance

Raster images have a fixed resolution, so they pixelate when scaled above their native size and waste storage when their content is piecewise-smooth (logos, icons, fonts, line art). Vector images sidestep both issues by describing the same content with curves and fills, which can be rasterized at any target resolution and are typically orders of magnitude smaller on disk. Automatic *vectorization* of an existing raster — recovering the underlying curves from a pixel grid — is therefore a recurring need in graphic design, OCR, font reconstruction, and low-bandwidth graphics.

2.2 Formal Definition

Input	$I : \{0, \dots, W-1\} \times \{0, \dots, H-1\} \rightarrow \{0, 1\}$, a binary raster image obtained by thresholding the luminance of an arbitrary PNG/JPG/BMP source.
Output	A scalable vector graphic $G \subset \mathbb{R}^2$, expressed as a union of closed paths each given by an outer ring (clockwise) and zero or more interior holes (counter-clockwise), filled with the even-odd rule.
Constraints	For a user-supplied tolerance $\delta > 0$ pixels, every pixel labelled foreground in I must lie inside G and every pixel labelled background must lie outside G , up to a Hausdorff distance of at most δ at the boundary.

2.3 Challenges

- **Complexity.** Naive pixel-outline emission is $\Theta(\partial G)$ in the boundary length, which can reach 10^5 vertices on a 25-megapixel input; we want output linear in the geometric complexity of the underlying shape, not in raster resolution.
- **Numerical instability.** Cubic-Bézier fitting via Bernstein polynomials and Newton-Raphson reparameterization is prone to round-off error; we therefore evaluate every curve with the de Casteljau subdivision scheme, which is provably backward-stable [4].
- **Degeneracies.** Closed polylines have coincident endpoints (RDP would collapse the loop); regions can have holes, figure-of-eight topology, or single-pixel “hairs” (which break Jacob’s stopping criterion if not handled).

3 Related Work and Baseline

The vectorization literature splits cleanly into the three stages we re-use here: contour extraction, polyline simplification, and curve fitting.

We define two *baselines* that share the same Moore-Neighbour front-end as the proposed pipeline but stop earlier:

Baseline-0 (raw). Emit the raw pixel-outline polygon directly as an SVG `<path>` consisting only of M/L/Z commands — no simplification, no fitting. This is the strawman a non-geometric implementation would produce.

Baseline-1 (polyline). Simplify each contour with RDP at tolerance ε , then emit the result as an M/L/Z polygon. This is what every classical CG textbook would recommend stopping at.

Existing vectorization tools. The dominant open-source vectorizer is POTRACE [6], which takes a binary bitmap and outputs smooth PostScript or SVG paths.

Table 1: Algorithms covered in class and their limitations.

Method	Time	Limitation
Pixel-outline polygon	$O(\partial G)$	Output size scales with raster, not geometry. No smoothness.
Moore-Neighbour tracing [5]	$O(\partial G)$	Cell-square boundaries; needs Jacob’s rule on figure-of-eight.
Ramer-Douglas-Peucker [2]	$O(n \log n)$ avg	Worst case $O(n^2)$; preserves vertices but not curvature.
de Casteljau [3, 4]	$O(d^2)$ per eval	Numerically stable but per-evaluation cost is $O(d^2)$.
Schneider fit [1]	$O(n \log n)$ avg	Greedy; requires reasonable endpoint tangent estimates.

POTRACE uses an integer-arithmetic polygon tracer comparable in spirit to Moore-Neighbour, a globally-optimal decomposition of the polygon into curve segments via a dynamic-programming sweep related to the Imai-Iri algorithm, and a final L^∞ -optimal cubic-Bézier fitting step. Its output quality is excellent, but the algorithm is tightly integrated and deliberately not decomposable into independent stages. By contrast, our pipeline is explicitly staged (Moore-Neighbour $\rightarrow$ RDP $\rightarrow$ Schneider), with each stage independently swappable. This staged design is what enables the ablation study below and why each stage maps cleanly to an algorithm analysed in lectures.

Inkscape’s *Trace Bitmap* dialog wraps a modified AUTOTRACE engine followed by optional POTRACE-style smoothing, providing a GUI at the cost of a large dependency footprint. Adobe Illustrator’s *Image Trace* extends this further with colour quantisation and stroke-width estimation. Both tools target the richer problem of full-colour vectorization and are therefore not directly comparable to our binary-only pipeline, but they confirm that the Moore-Neighbour $\rightarrow$ curve-fit paradigm is the production standard.

4 Proposed Approach

4.1 Intuition

Both baselines describe the boundary as a finite set of *linear* segments. But the underlying shape we are trying to recover is typically smooth (logos, fonts, silhouettes), so a *cubic* fit should be able to capture the same boundary with dramatically fewer primitives, at the cost of a tiny per-primitive overhead. The proposed pipeline therefore feeds the RDP output into Schneider’s recursive least-squares cubic Bézier fitter.




Figure 1: The proposed pipeline. Baseline-0 stops after the trace stage; baseline-1 stops after RDP.

4.2 Algorithm Design

Stage 1 – Binarization. Each input pixel is converted to luminance $Y = 0.299R + 0.587G + 0.114B$ (Rec.601), composited over white if it has alpha, and classified as foreground iff $Y < \tau$ for a user-controlled threshold $\tau \in [0, 255]$.

Stage 2 – Connected components and Moore-Neighbour tracing. Both foreground and background are flood-filled with 4-connectivity. The unique background component touching the image border is the *exterior*; every other background component is a *hole* of the foreground component that surrounds it. For each component R , we start at its topmost-leftmost pixel p_0 (so the cell to its west is guaranteed outside R) and repeatedly: (i) rotate clockwise through the eight neighbours of the current boundary pixel starting just after the previous pixel; (ii) the first neighbour belonging to R becomes the new boundary pixel. Termination uses Jacob’s criterion — re-entering p_0 from the same previous cell as at the start — which avoids spurious early termination on figure-of-eight shapes.

Stage 3 – Ramer–Douglas–Peucker.

```

RDP(P[i..j], eps):
    dmax = 0; k = i
    for m = i+1 .. j-1:
        d = perp_distance(P[m], line(P[i], P[j]))
        if d > dmax: dmax = d; k = m
    if dmax > eps:
        RDP(P[i..k], eps); RDP(P[k..j], eps)
    else:
        emit segment (P[i], P[j])
  
```

The perpendicular distance from p to the line through a, b is

$$\text{dist}(p, \overline{ab}) = \frac{|(b_x - a_x)(p_y - a_y) - (b_y - a_y)(p_x - a_x)|}{\|b - a\|}.$$

Closed polylines are split first at the vertex farthest from p_0 to avoid collapsing the loop.

Stage 4 – de Casteljau and Schneider’s fit. A cubic Bézier $\mathbf{B}(t) = \sum_{i=0}^3 \binom{3}{i} (1-t)^{3-i} t^i \mathbf{V}_i$ is evaluated by

$$\mathbf{V}_i^{(r)} = (1-t)\mathbf{V}_i^{(r-1)} + t\mathbf{V}_{i+1}^{(r-1)}, \quad r = 1, 2, 3, \quad (1)$$

with $\mathbf{B}(t) = \mathbf{V}_0^{(3)}$. The recurrence reaches backward error $O(\epsilon_{\text{mach}})$ regardless of t , whereas direct Bernstein evaluation can lose half its bits when t is close to either endpoint [4]. The same recurrence applied to the quadratic control points $\mathbf{Q}_i = 3(\mathbf{V}_{i+1} - \mathbf{V}_i)$ yields the first derivative, which we need for the Newton refinement below. Schneider’s recursive fit is then:

```

FitCurve(P[0..n], delta):
    t1 = unit(P[1] - P[0])           # left tangent
    t2 = unit(P[n-1] - P[n])         # right tangent
    return FitOne(P, 0, n, t1, t2, delta)

FitOne(P, i, j, t1, t2, delta):
    if j - i == 1:
        return [chord-segment(P[i], P[j])]
    u = ChordLengthParam(P, i, j)    # u_i in [0,1]
    bez = SolveLeastSquares(P, i, j, u, t1, t2)
    (err, k) = MaxError(P, i, j, u, bez)
    if err < delta^2:
        return [bez]
    if err < (delta * iterErrFactor)^2:
        u' = NewtonStep(P, i, j, u, bez)
        bez' = SolveLeastSquares(P, i, j, u', t1, t2)
        if MaxError(...) < delta^2: return [bez']
    # otherwise: subdivide at the worst-error vertex
    c = unit((P[k+1] - P[k-1]) / 2)
    L = FitOne(P, i, k, t1, c, delta)
    R = FitOne(P, k, j, -c, t2, delta)
    return L ++ R
  
```

The least-squares step solves a 2×2 system for the two tangent magnitudes α_1, α_2 that minimise $\sum_i \|\mathbf{B}(u_i) - \mathbf{d}_i\|^2$ subject to fixed endpoints $\mathbf{V}_0, \mathbf{V}_3$ and fixed tangent directions $\hat{\mathbf{t}}_1, \hat{\mathbf{t}}_2$ — a closed-form solve that reduces to a 2×2 matrix inversion.

4.3 Theoretical Analysis

Let $A = WH$ be the image area and n the total raw vertex count returned by Moore-Neighbour. Then $n \leq 4(W + H)$ in the worst case and is typically $\ll A$. The per-stage costs are

Stage	Time	Space
Binarize+CC label	$O(A)$	$O(A)$
Moore-Neighbour	$O(n)$	$O(1)$
RDP	$O(n \log n)$ avg / $O(n^2)$ wc	$O(n)$
Schneider fit	$O(n \log n)$ avg	$O(n)$
SVG emit	$O(\text{out})$	$O(\text{out})$

Total: $O(A + n \log n)$ time, dominated by the connected-component flood-fill at large image area and by RDP at small.

5 Implementation Details

Item	Detail
Language	C++17, no exceptions outside I/O.
Build system	CMake ≥ 3.15 (Apple Clang 17 verified).
Dependencies	<code>stb_image.h</code> (public domain) only.
Data structs	Row-major byte raster for the image; <code>std::vector<Point></code> for polylines; an <code>SvgPath</code> carrying one outer ring + k holes for the SVG emitter.
Edge cases	Closed-loop RDP split (§4), hole/parent assignment via background flood-fill, alpha compositing over white before binarization, Jacob’s criterion in the tracer, FP precision handled by working in pixel-centre coordinates.

The de Casteljau evaluator is the kernel that ties everything together:

```
Point deCasteljau(const std::array<Point,4>& V,
double t) {
    std::array<Point,4> q = V;
    for (int level = 1; level < 4; ++level)
        for (int i = 0; i < 4 - level; ++i)
            q[i] = lerp(q[i], q[i+1], t);
    return q[0];
}
```

Module Structure

Table 2: Source-level module breakdown (LOC = lines of C++ without blank lines or comments).

Module	Responsibility	LOC
<code>image</code>	Load, binarize, flood-fill, CC label	220
<code>tracer</code>	Moore-Neighbour + Jacob’s criterion	130
<code>rdp</code>	RDP, closed-loop split	90
<code>bezier</code>	de Casteljau evaluator, Schneider fit	280
<code>svg</code>	SVG path builder and file emitter	110
<code>main</code>	CLI, option parsing, per-stage timing	120
<code>gen_synthetic</code>	Synthetic image generator	60
Total		≈ 1010

Command-line interface. The binary `build/vectorizer` exposes all three pipeline modes through a single flag:

```
vectorizer <input> -o <output.svg>
--mode raw|polyline|bezier (default: bezier)
--eps <rdp_tol_px> (default: 1.5)
--delta <bez_tol_px> (default: 2.0)
--threshold <0-255> (default: 128)
--bench # emit CSV timing
line
```

All three modes are selectable at runtime without re-compilation, which makes the three-way comparison in

Section 7 clean: the same binary is invoked three times on the same input with only the `--mode` flag changed.

SVG output format. Each connected component is emitted as a single `<path>` element. The outer ring uses clockwise winding; interior holes use counter-clockwise winding, separated by `ZM` tokens within the same path. The `fill-rule="evenodd"` attribute renders holes correctly without explicit Boolean subtraction. In `bezier` mode the path uses `MCZ` commands; in `polyline` mode, `MLZ`. Coordinates are written as integer pixel values, keeping file size minimal and avoiding floating-point formatting overhead.

6 Correctness and Validation

Smoke-test suite. The script `tests/test_algorithms.sh` checks three pipeline properties on a fixed synthetic image (100×100 px, 33-px radius disk with a concentric hole):

Raw mode — vertex count. The Moore-Neighbour boundary must contain at least 400 vertices ($2\pi r \approx 207$ for a single disk; the 8-connected Moore-Neighbour walk yields roughly twice that because it traces the cell boundary, not the pixel centre path). The SVG byte count is checked to be consistent with this count, catching any silent truncation in the emitter.

Polyline mode — RDP compression. After RDP at $\varepsilon = 1.0$ px the simplified vertex count must fall below 200 (roughly a 4× reduction on a circle, consistent with the $O(\sqrt{n})$ worst-case guarantee for convex curves under RDP).

Bézier mode — cubic economy. The Schneider fit must produce fewer than 30 cubic Bézier segments (a unit circle needs as few as 4 segments at $\delta = 1$ px) and at least 2 (a single segment cannot close a loop without degeneracy).

Assertions. Each check is implemented with `assert_le` / `assert_ge` shell helpers that parse the CSV line emitted by `--bench --quiet` and compare the relevant column to expected thresholds. The script exits non-zero on the first failure, integrating directly with `make test`.

Determinism. The pipeline contains no random state; given the same binary and input, the output SVG is byte-identical across consecutive runs. This property is essential for the median benchmark calculation: the five repeated runs differ only in OS-level noise (cache warmth, scheduler jitter), not in computation.

Topology correctness. The disk image includes an interior hole (the lighter annulus centre), which exercises the full hole-detection path: the background

flood-fill must correctly classify the enclosed background region as a hole of the enclosing foreground ring, not as a separate outer boundary. The image-border-touching exterior background is correctly distinguished from the internal hole by the border-connectivity test in `image.cpp`. Together, these conditions exercise every branch of the topology-assignment code. The test asserts that the number of `<path>` elements in the output SVG equals the expected number of connected components for the given tile layout.

7 Experimental Setup

Datasets. Synthetic raster images are produced by `build/gen_synthetic`: each is a $g \times g$ tiled grid of identical concentric annuli (a filled disk with a smaller circular hole), so the total Moore-Neighbour boundary length is approximately $8rg^2$. By choosing $(r, g) \in \{(128, 1), (80, 4), (85, 12)\}$ we hit $n \approx 10^3, 10^4, 10^5$ raw boundary vertices respectively, with raster sizes of only $272^2, 704^2$ and 2232^2 pixels. The grid layout is what makes $n = 10^5$ feasible at all: a single annulus large enough to give 10^5 boundary pixels would require a $\sim 25\text{k} \times 25\text{k}$ raster ($\sim 650\text{ MB}$), exceeding the safety limits of `stb_image`.

Environment. Apple M2 (8 cores), 16 GB RAM, macOS Sequoia 15.4, Apple Clang 17 with `-O3`. Each (n, mode) cell is the median of five back-to-back runs.

Metrics. We report runtime (ms; total and per-stage), output SVG size in bytes, and number of geometric primitives in the output (vertices for polyline modes, cubic segments for Bézier mode). Accuracy is bounded by construction: every pipeline obeys the user-supplied tolerances ε (RDP, in pixels) and δ (Bézier, in pixels). We use $\varepsilon = \delta = 1$ pixel throughout.

Reproducibility. All numbers below are reproduced by `make bench` (which calls `experiments/benchmark.sh` and writes `experiments/results/runtime.csv`); plots are regenerated by `make plot`.

8 Results and Evaluation

8.1 Output Compactness

Table 3: Output size in bytes for the same input, three modes, three input scales. Reduction is relative to the **raw** baseline.

n	raw	polyline	bezier	reduction
10^3	16 783	1 053	619	27.1×
10^4	172 522	15 146	8 082	21.3×
10^5	1 754 478	144 798	84 630	20.7×

8.2 Primitive Count

Table 4: Number of geometric primitives in the output (vertices for polyline modes, cubic-Bézier segments for the proposed mode).

n (raw v.)	raw v.	polyline v.	bezier seg.
10^3	962	50	8
10^4	9 696	800	128
10^5	91 872	7 200	1 296

8.3 Runtime (Scaling Test)

Table 5: Median total runtime in ms over five runs. “Speedup” is the proposed pipeline’s runtime relative to the raw baseline.

n (raw v.)	raw (ms)	polyline (ms)	bezier (ms)	speedup
10^3	1.39	1.02	1.03	1.36×
10^4	17.25	13.90	13.83	1.25×
10^5	841.4	814.3	800.5	1.05×

Two log-log plots produced by `visualization/plot_results.py` show that all three pipelines are strictly linear in n in this range, with the constant for the proposed pipeline always below the constant for the baselines on both axes.




Figure 2: Median total runtime (log-log). All three pipelines scale linearly in n ; the proposed Bézier mode is consistently the fastest because its much smaller output dominates SVG serialization cost.

8.4 Discussion of Results

The proposed pipeline wins on both metrics simultaneously, which is *not* obvious a priori — one would expect curve fitting to add overhead. Two effects explain the result.

First, the dominant runtime cost is Moore-Neighbour tracing ($>95\%$ of total at $n = 10^5$), which is shared across all three modes. The fit step itself contributes only 0.63 ms even on the $n = 10^5$ input.




Figure 3: Output SVG size (log-log). Both baselines grow linearly in n ; the proposed pipeline grows linearly as well but with a constant that is 20–27 $\times$ smaller than the raw baseline.

Second, the SVG *serializer* processes far less data when the output is sparser. At $n = 10^5$ the proposed mode has to write 84kB of text instead of 1.7MB, which more than compensates for the 0.63 ms of fitting. This is why the proposed pipeline is the fastest of the three at every scale, despite doing strictly more work upstream.

The compression ratio is roughly constant at $\sim 21\times$ once $n \geq 10^4$, which matches the expected behaviour: each cubic Bézier segment captures a region of bounded curvature and therefore absorbs an essentially constant number of RDP vertices, regardless of n .

8.5 Per-Stage Runtime Breakdown

To understand *why* the proposed pipeline is the fastest of the three despite doing strictly more work, Table 6 breaks the median runtime at $n = 10^5$ down by stage. The pattern is the same at smaller n , just scaled.

Table 6: Per-stage runtime in ms at $n = 10^5$. “serialize” is the total minus the sum of the named stages and absorbs SVG-emission cost, which scales with the output size.

Stage	raw	polyline	bezier
load+binarize	13.99	13.79	13.48
trace	790.5	796.0	783.0
RDP simplify	0	1.27	1.16
Bézier fit	0	0	0.63
serialize	36.9	3.2	2.3
total	841.4	814.3	800.5

Two observations follow. (i) Tracing alone is responsible for $\sim 98\%$ of total runtime in every mode — the cost the three modes share dwarfs everything they differ on. (ii) The **raw** mode pays a 36.9ms serialization tax because it has to write out 1.7MB of L-commands; the proposed mode writes only 84kB and pays just 2.3ms. This 34.6ms of saved serialization more than

offsets the 1.79ms that RDP+Schneider together cost, delivering a net 5% wall-clock win at the largest tested scale.

8.6 Visual Output

A qualitative side-by-side comparison is included in the project repository as `examples/disk.png` (input raster, 200×200) and `examples/disk.svg` (proposed-mode output). Opening the SVG in any browser and varying the zoom level visually confirms the two key properties promised by vectorization: the curve remains crisp at any rendering size, and the file is one or two orders of magnitude smaller than even a modestly compressed PNG of the same shape. The qualitative result on the annulus is a single closed loop of nine cubic-Bézier segments that is visually indistinguishable from a perfect circle.

8.7 Worst Case

The Bézier fitter degenerates back to short segments on highly non-smooth boundaries (e.g. rasterized text at small sizes), where the recursive subdivision in Schneider’s algorithm bottoms out. We verified that on such inputs the segment count approaches the RDP vertex count and the size advantage shrinks; runtime, however, is still bounded by tracing.

9 Discussion

Strengths. Three textbook algorithms compose, in ~ 1 k lines of C++ and zero external runtime dependencies, into a vectorizer that beats both polygonal baselines on *both* runtime and output size simultaneously. Each stage is independently swappable (polyline-only mode is a one-flag change), which made the experiment easy to implement.

Limitations. (i) The current binarizer assumes a single foreground colour; photographs need a colour-quantization pre-pass. (ii) The Moore-Neighbour tracer treats each pixel as a square sample; marching squares would yield sub-pixel-accurate seeds at no extra algorithmic cost. (iii) Schneider’s fit is greedy — a globally optimal L^∞ segmentation à la Imai-Iri would tighten the segment count further but with diminishing returns for typical artwork.

Trade-offs. Lowering the user tolerances (ϵ, δ) moves the pipeline along a smooth Pareto front: smaller output at higher fidelity costs strictly more segments and time. The defaults (1.5, 2.0) px were chosen to be visually lossless on ~ 1 megapixel inputs.

10 Conclusion

We implemented and benchmarked a complete raster-to-vector pipeline that composes Moore-Neighbour

boundary tracing, Ramer–Douglas–Peucker simplification, and Schneider’s de-Casteljau-based cubic Bézier fit. Against two polygonal baselines on synthetic inputs spanning three orders of magnitude in boundary size, the proposed pipeline produces SVG outputs that are 20–27× smaller than the raw baseline and 1.7–1.8× smaller than the RDP-only baseline, while remaining the fastest end-to-end at every scale tested. The result is a self-contained portfolio-quality tool (< 1000 LOC, one `cmake --build`) that directly demonstrates the geometric trade-offs taught in the course.

Future work includes colour-quantized multi-layer vectorization, sub-pixel marching-squares seeds, and a corner-detection pre-pass to preserve sharp angles exactly. All source code, benchmarks, and the demo CLI are available in the project repository.

References

- [1] P. J. Schneider, “An Algorithm for Automatically Fitting Digitized Curves,” in *Graphics Gems*, A. S. Glassner, Ed., Academic Press, 1990, pp. 612–626.
- [2] D. H. Douglas and T. K. Peucker, “Algorithms for the reduction of the number of points required to represent a digitized line or its caricature,” *Cartographica*, vol. 10, no. 2, pp. 112–122, 1973.
- [3] P. de Casteljau, *Outillage Méthodes Calcul*, Tech. Rep., André Citroën Automobiles, 1959.
- [4] R. T. Farouki, “The Bernstein polynomial basis: A centennial retrospective,” *Computer Aided Geometric Design*, vol. 29, no. 6, pp. 379–419, 2012.
- [5] R. C. Gonzalez and R. E. Woods, *Digital Image Processing*, 4th ed., Pearson, 2018, ch. 11 (Boundary tracing).
- [6] P. Selinger, “Potrace: A polygon-based tracing algorithm,” *Proceedings of the Free Software and Open Source Symposium*, 2003. Available: <https://potrace.sourceforge.net/potrace.pdf>.

Appendix A: Reproducing the Numbers in This Report

From a clean checkout:

```
cmake -S . -B build -DCMAKE_BUILD_TYPE=Release
cmake --build build -j
make bench # writes experiments/results/runtime.csv
make plot # writes experiments/results/*.png
make test # runs tests/test_algorithms.sh
make demo # runs the CLI demo
```

The CSV columns are `n, mode, W, H, raw_v, simp_v, out_seg, out_bytes, t_load, t_trace, t_simp, t_fit, t_total`.

Appendix B: Derivation of the 2 × 2 Least-Squares System

The Schneider fit reduces curve fitting to a small linear system at each recursive call. We derive it here for completeness.

Setup. Given $m+1$ data points $\mathbf{d}_0, \dots, \mathbf{d}_m$ with fixed endpoints $\mathbf{d}_0 = \mathbf{V}_0$, $\mathbf{d}_m = \mathbf{V}_3$, and chord-length parameters

$$u_i = \frac{\sum_{k=1}^i \|\mathbf{d}_k - \mathbf{d}_{k-1}\|}{\sum_{k=1}^m \|\mathbf{d}_k - \mathbf{d}_{k-1}\|},$$

we parameterise the two free interior control points as $\mathbf{V}_1 = \mathbf{V}_0 + \alpha_1 \hat{\mathbf{t}}_1$ and $\mathbf{V}_2 = \mathbf{V}_3 + \alpha_2 \hat{\mathbf{t}}_2$ for scalars $\alpha_1, \alpha_2 \geq 0$ and unit endpoint tangents $\hat{\mathbf{t}}_1, \hat{\mathbf{t}}_2$.

Objective. Minimise the squared residual $E = \sum_i \|\mathbf{B}(u_i) - \mathbf{d}_i\|^2$. Expanding the cubic Bernstein form $\mathbf{B}(u) = \sum_{k=0}^3 B_k(u) \mathbf{V}_k$ and substituting the parameterised control points,

$$\mathbf{B}(u_i) = \alpha_1 B_1(u_i) \hat{\mathbf{t}}_1 + \alpha_2 B_2(u_i) \hat{\mathbf{t}}_2 + \underbrace{B_0(u_i) \mathbf{V}_0 + B_3(u_i) \mathbf{V}_3}_{\mathbf{C}_i}.$$

Define $\mathbf{r}_i = \mathbf{d}_i - \mathbf{C}_i$. Setting $\partial E / \partial \alpha_k = 0$ yields the 2×2 normal system

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{12} & C_{22} \end{pmatrix} \begin{pmatrix} \alpha_1 \\ \alpha_2 \end{pmatrix} = \begin{pmatrix} X_1 \\ X_2 \end{pmatrix}, \quad (2)$$

with coefficients

$$\begin{aligned} C_{11} &= \sum_i B_1(u_i)^2, \\ C_{22} &= \sum_i B_2(u_i)^2, \\ C_{12} &= \sum_i B_1(u_i) B_2(u_i) (\hat{\mathbf{t}}_1 \cdot \hat{\mathbf{t}}_2), \\ X_k &= \sum_i B_k(u_i) (\mathbf{r}_i \cdot \hat{\mathbf{t}}_k), \quad k = 1, 2. \end{aligned}$$

Stability. The determinant $\Delta = C_{11}C_{22} - C_{12}^2 \geq 0$ by Cauchy–Schwarz. When Δ is below machine precision relative to $\max(C_{11}, C_{22})$ (nearly-parallel tangents or degenerate chord length), the system is ill-conditioned and the implementation falls back to the heuristic $\alpha_1 = \alpha_2 = \frac{1}{3} \|\mathbf{V}_3 - \mathbf{V}_0\|$, which places interior control points one-third of the chord length from each endpoint — the well-known cubic arc approximation starting point.

Complexity. Each evaluation of (2) costs $O(m)$ dot products with $O(1)$ Bernstein evaluations per point. Across all recursive segments the total fitting work is $O(n)$ per polyline, matching the RDP stage and leaving Moore-Neighbour tracing as the sole asymptotic bottleneck.